

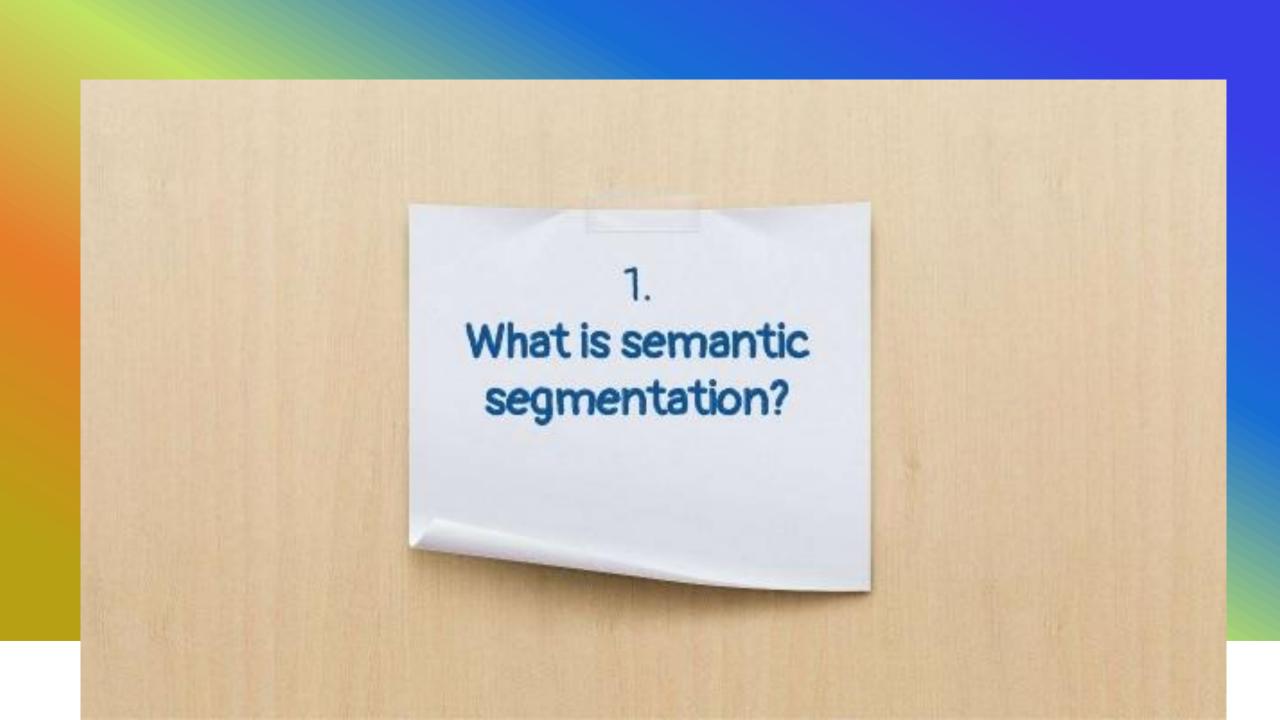
Semantic Segmentation

Deep Learning based Semantic Segmentation using Drone Image Dataset

Anam Nawaz Khan MCL AD20206809

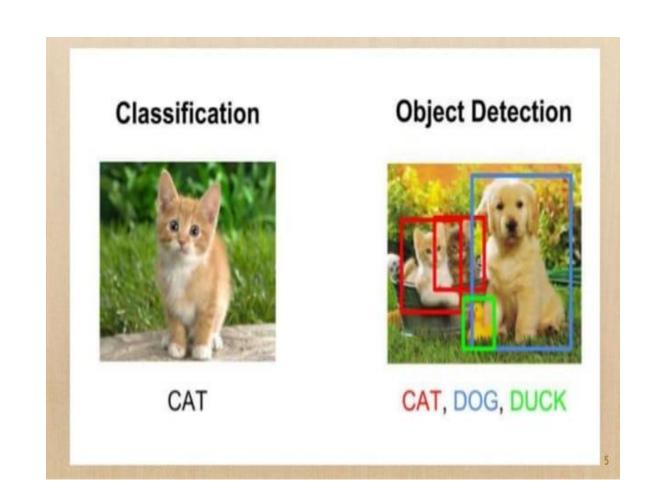
Table of Contents:

- + Introduction
- + Applications
- + Representations
- + Architecture Construction
- + Implementation and Results

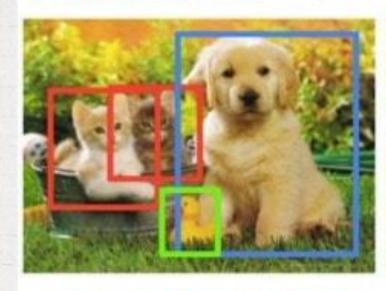


Limitations

Still a bit rough since we are only drawing bounding boxes and don't get an accurate idea of object shape



What if!?



CAT, DOG, DUCK



CAT, DOG, DUCK

What is semantic segmentation

- + Semantic image segmentation is the task of classifying each pixel in an image from a predefined set of classes. In the following example, different entities are classified.
- + In the above example, the pixels belonging to the bed are classified in the class "bed", the pixels corresponding to the walls are labeled as "wall", etc.
- + In particular, our goal is to take an image of size W x H x 3 and generate a W x H matrix containing the predicted class ID's corresponding to all the pixels.
- + Therefore In order to perform semantic segmentation, a higher level understanding of the image is required.
- + The algorithm should figure out the objects present and also the pixels which correspond to the object.
- + Semantic segmentation is one of the essential tasks for complete scene understanding.



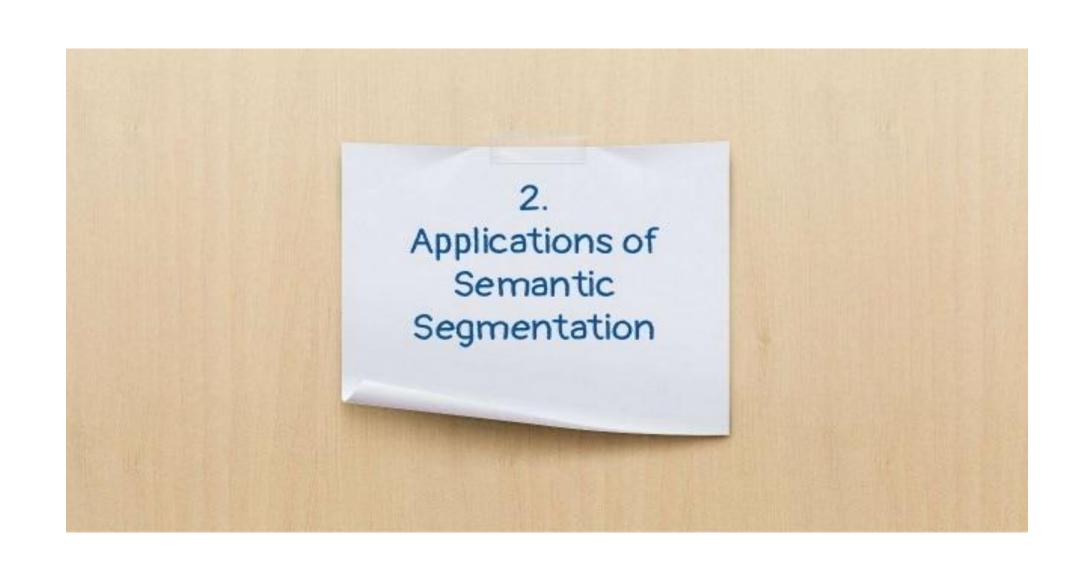






1: Person 2: Purse 3: Plants/Grass 4: Sidewalk

Input



Semantic Segmentation Applications

There are several applications for which semantic segmentation is very useful.

Medical images

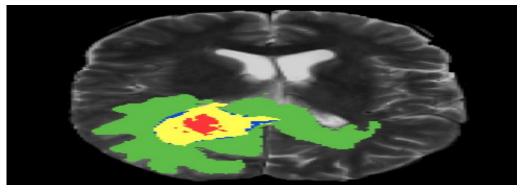
+ Automated segmentation of body scans can help doctors to perform diagnostic tests. For example, models can be trained to segment tumor.

Autonomous vehicles

+ Autonomous vehicles such as self-driving cars and drones can benefit from automated segmentation. For example, self-driving cars can detect drivable regions.

Satellite image analysis

+ Aerial images can be used to segment different types of land. Automated land mapping can also be done.

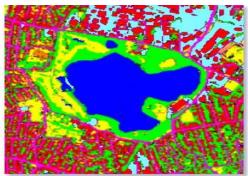


Tumor segmentation of brain MRI scan



Segmentation of a road scene

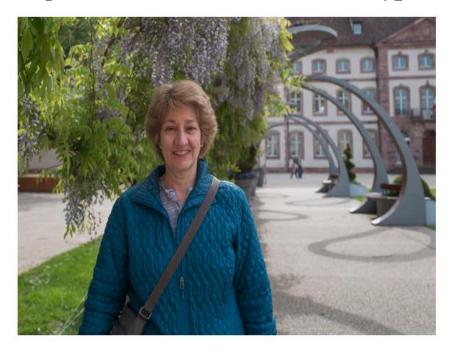




Segmentation of a satellite image



Our goal is to take either a RGB color image or a gray-scale image and output a segmented map where each pixel contain a class label represented as integer





- 1: Person
- 2: Purse
- 3: Plants/Grass
- 4: Sidewall
- 5: Building/Structures

```
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
      3
```

Input Semantic Labels



Aerial Semantic Segmentation Drone Dataset

Dataset Overview

- + The Semantic Drone Dataset focuses on semantic understanding of urban scenes for increasing the safety of autonomous drone flight and landing procedures.
- + The imagery depicts more than 20 houses from bird's eye view acquired at an altitude of 5 to 30 meters above ground.
- + A high resolution camera was used to acquire images at a size of 6000x4000px (24Mpx).
- + The training set contains 400 publicly available images and the test set is made up of 200 private images.
- + We prepared pixel-accurate annotation for the same training and test set. The complexity of the dataset is limited to 20 classes as listed in the following table.

Table 1: Semantic classes of the Drone Dataset

tree	rocks	dog	fence
gras	water	car	fence-pole
other vegetation	paved area	bicycle	window
dirt	pool	roof	door
gravel	person	wall	obstacle





Libraries

Importing the Necessary Libraries

- + NumPy is a Python library used for working with arrays.
- + Pandas is a predominantly used python data analysis library.
- + Matplotlib is a cross-platform, data visualization and graphical plotting library for Python and its numerical extension NumPy.
- + Imageio is a Python library that provides an easy interface to read and write a wide range of image data
- + Input() is used to instantiate a Keras tensor. A Keras tensor is a symbolic tensor-like object, which we augment with certain attributes that allow us to build a Keras model by knowing the inputs and outputs of the model.
- + Keras Conv2D is a 2D Convolution Layer, this layer creates a convolution kernel that is wind with layers input which helps produce a tensor of outputs.
- + MaxPooling2D layer Down-samples the input along its spatial dimensions (height and width) by taking the maximum value over an input window (of size defined by pool size) for each channel of the input.
- + Dropout: A Simple Way to Prevent Neural Networks from Overfitting
- + Conv2DTranspose Using Conv2DTranspose will up-sample its input but the key difference is the model should learn what is the best up-sampling for the job. conv2d_transpose() simply transposes the weights and flips them by 180 degrees. Then it applies the standard conv2d(). "Transposes" practically means that it changes the order of the "columns" in the weights tensor.
- + Concatenate is a Layer that concatenates a list of inputs. It takes as input a list of tensors, all of the same shape except for the concatenation axis, and returns a single tensor that is the concatenation of all inputs.

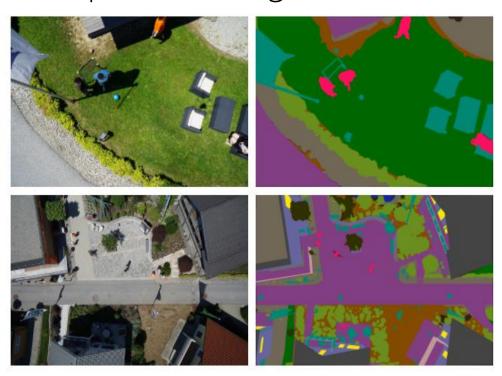
```
import os
import numpy as np
import pandas as pd
import imageio
import matplotlib.pyplot as plt
%matplotlib inline
```

```
import tensorflow as tf
import numpy as np

from tensorflow.keras.layers import Input
from tensorflow.keras.layers import Conv2D
from tensorflow.keras.layers import MaxPooling2D
from tensorflow.keras.layers import Dropout
from tensorflow.keras.layers import Conv2DTranspose
from tensorflow.keras.layers import concatenate
```

Load and Split The Data

+ The training set contains 400 publicly available images and the test set is made up of 200 private images.



```
import os
import numpy as np
import pandas as pd
import imageio
import matplotlib.pyplot as plt
%matplotlib inline
image_path = os.path.join(path,'../input/semantic-drone-dataset/dataset/semantic_drone_dataset/original_images/')
mask_path = os.path.join(path,'../input/semantic-drone-dataset/dataset/semantic_drone_dataset/label_images_semantic/')
image_list = os.listdir(image_path)
mask_list = os.listdir(mask_path)
image_list = [image_path+i for i in image_list]
mask_list = [mask_path+i for i in mask_list]
image_list = sorted(image_list)
mask_list = sorted(mask_list)
print("number of images is : {} ".format(len(image_list)))
print(image_list[0])
print(mask_list[0])
```

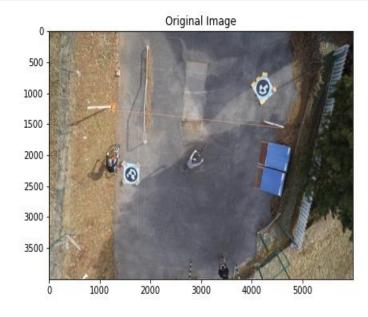
```
number of images is : 400
../input/semantic-drone-dataset/dataset/semantic_drone_dataset/original_images/000.jpg
../input/semantic-drone-dataset/dataset/semantic_drone_dataset/label_images_semantic/000.png
```

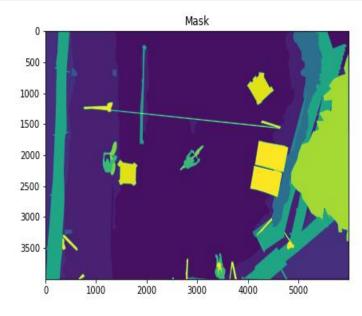
Exploration and overview of images

Drone's images semantic segmentation overview

```
n = 10 # you can chose any index
img = imageio.imread(image_list[n])
print(img.shape)
mask = imageio.imread(mask_list[n])
print(mask.shape)

# now let's plot
fig ,arr = plt.subplots(1,2,figsize=(15,10))
arr[0].imshow(img)
arr[0].set_title('Original Image')
arr[1].imshow(mask)
arr[1].set_title('Mask')
```





Preprocessing our data

- + Tf.io.read_file

 Returns the contents of the file
- + Tf.img.decode_png

 Decodes an encoded image to a unit8 or uint16 tensor
- + Tf.image.convert_image_dtype Supported datatype of the image
- + Tf.reduce_max

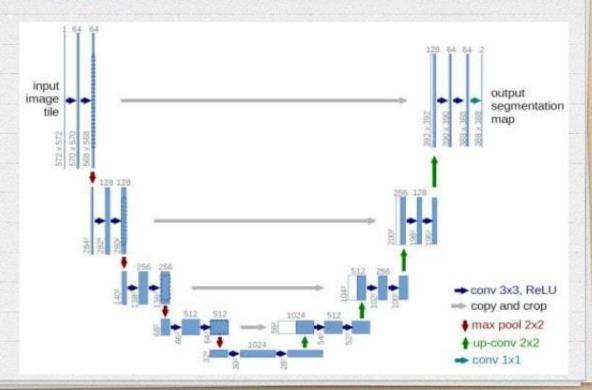
 Reduce the dimensions along
 the dimensions given in axis
- + Tf.image.resize

 Resize the images using a specified method

```
def process_path(image_path, mask_path):
    img = tf.io.read_file(image_path)
    img = tf.image.decode_png(img,channels=3)
    img = tf.image_convert_image_dtype(img,tf.float32) #this do the same as dividing by 255 to set the values between \theta and 1 (normalization)
    mask = tf.io.read_file(mask_path)
    mask = tf.image.decode_png(mask,channels=3)
    mask = tf.math.reduce_max(mask,axis=-1,keepdims=True)
    return img , mask
def preprocess(image, mask) :
    input_image = tf.image.resize(image, (96, 128), method='nearest')
    input_mask = tf.image.resize(mask,(96,128),method='nearest')
    return input_image , input_mask
image_ds = dataset.map(process_path) # apply the preproces_path function to our dataset
print(image_ds)
processed_image_ds = image_ds.map(preprocess) # apply the preprocess function to our dataset
```

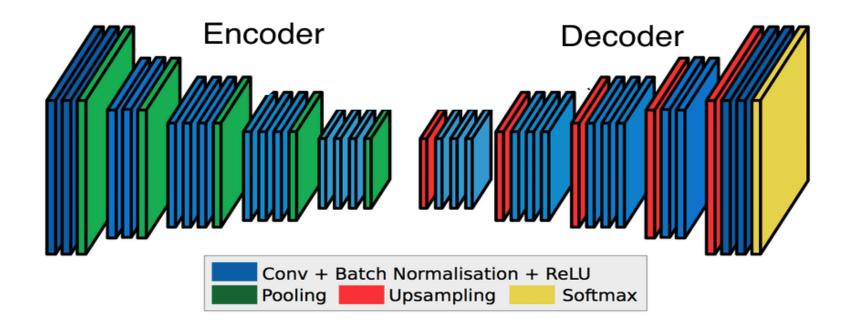
U-Net Architecture..

Consists of a contracting path to capture context and a symmetric expanding path that enables precise localization.



Encoder-decoder Structure

- + This is called an encoder-decoder structure. Where the layers which down-sample the input are the part of the encoder and the layers which up-sample are part of the decoder
- + When the model is trained for the task of semantic segmentation, the encoder outputs a tensor containing information about the objects, and its shape and size. The decoder takes this information and produces the segmentation maps.



Define The Conv Block For The Contracting Path

- + The encoder network acts as the feature extractor and learns an abstract representation of the input image through a sequence of the encoder blocks. Here number of filters are 32 and two consecutive Convolution Layers are applied
- + In the Encoder, the size of the image gradually reduces while the depth gradually increases.
- + Each encoder block consists of two 3×3 convolutions, where each convolution is followed by a ReLU (Rectified Linear Unit) activation function
- + The ReLU activation function introduces non-linearity into the network, which helps in the better generalization of the training data.
- + The block gives two result values as a return value. The first value is the activation value to be used in the next layer, and the second value is the value before pooling for the skip connection. The initial value declaration of the filter uses the He-Initialization method
- + The output of the ReLU acts as a skip connection for the corresponding decoder block. These skip connections provide additional information that helps the decoder to generate better semantic features
- + Next, follows a 2×2 max-pooling, where the spatial dimensions (height and width) of the feature maps are reduced by half.
- + This reduces the computational cost by decreasing the number of trainable parameters.

```
def conv_block(inputs=None, n_filters=32, dropout_prob=0, max_pooling=True):
   conv = Conv2D(n_filters.
                  kernel_size = 3.
                  activation='relu'.
                  padding='same',
                  kernel_initializer=tf.keras.initializers.HeNormal())(inputs)
    conv = Conv2D(n_filters,
                  kernel size = 3.
                  activation='relu'.
                  padding='same',
                  kernel_initializer=tf.keras.initializers.HeNormal())(conv)
   if dropout_prob > 0:
       conv = Dropout(dropout_prob)(conv)
   if max_pooling:
       next_layer = MaxPooling2D(pool_size=(2,2))(conv)
    else:
       next_layer = conv
    skip_connection = conv
    return next_layer, skip_connection
```

Define the up-sampling block for the expanding path

- + The expansion path (Decoder) where we apply transposed convolutions along with regular convolutions
- + In the decoder, the size of the image gradually increases and the depth gradually decreases.
- + The decoder network is used to take the abstract representation and generate a semantic segmentation mask.
- + Intuitively, the Decoder recovers the "WHERE" information (precise localization) by gradually applying up-sampling
- + To get better precise locations, at every step of the decoder we use skip connections by concatenating the output of the transposed convolution layers with the feature maps from the Encoder at the same level:
- + After every concatenation we again apply two consecutive regular convolutions so that the model can learn to assemble a more precise output
- + This is what gives the architecture a symmetric U-shape, hence the name UNET
- + On a high level, we have the following relationship: Input (128x128x1) (spatial dimx feature channels)=> Encoder =>(8x8x256) => Decoder => Ouput (128x128x1)
- + The U-Net combines the location information from the down sampling path to finally obtain a general information combining localization and context, which is necessary to predict a good segmentation map.

```
def upsampling_block(expansive_input, contractive_input, n_filters=32):
   up = Conv2DTranspose(
                n_filters,
                kernel_size = 3.
                strides=(2,2),
                padding='same')(expansive_input)
   merge = concatenate([up, contractive_input], axis=3)
   conv = Conv2D(n_filters,
                kernel_size = 3.
                activation='relu',
                padding='same',
                kernel_initializer=tf.keras.initializers.HeNormal())(merge)
   conv = Conv2D(n_filters,
                kernel_size = 3,
                activation='relu'.
                padding='same',
                kernel_initializer=tf.keras.initializers.HeNormal())(conv)
   return conv
```

Define the u-net model

 we will Define the unet model, which composes of a set of conv blocks and upsampling blocks

concatenate (Concatenate)	(None,	12,	16,	512)	0	conv2d_transpose[0][0] dropout[0][0]
conv2d_10 (Conv2D)	(None,	12,	16,	256)	1179904	concatenate[0][0]
conv2d_11 (Conv2D)	(None,	12,	16,	256)	590080	conv2d_10[0][0]
onv2d_transpose_1 (Conv2DTrans	(None,	24,	32,	128)	295040	conv2d_11[0][0]
concatenate_1 (Concatenate)	(None,	24,	32,	256)	0	conv2d_transpose_1[0][0] conv2d_5[0][0]
conv2d_12 (Conv2D)	(None,	24,	32,	128)	295040	concatenate_1[0][0]
conv2d_13 (Conv2D)	(None,	24,	32,	128)	147584	conv2d_12[0][0]
conv2d_transpose_2 (Conv2DTrans	(None,	48,	64,	64)	73792	conv2d_13[0][0]
concatenate_2 (Concatenate)	(None,	48,	64,	128)	0	conv2d_transpose_2[0][0] conv2d_3[0][0]
conv2d_14 (Conv2D)	(None,	48,	64,	64)	73792	concatenate_2[0][0]
conv2d_15 (Conv2D)	(None,	48,	64,	64)	36928	conv2d_14[0][0]
conv2d_transpose_3 (Conv2DTrans	(None,	96,	128	, 32)	18464	conv2d_15[0][0]
concatenate_3 (Concatenate)	(None,	96,	128	, 64)	0	conv2d_transpose_3[0][0] conv2d_1[0][0]
conv2d_16 (Conv2D)	(None,	96,	128	, 32)	18464	concatenate_3[0][0]
onv2d_17 (Conv2D)	(None,	96,	128	, 32)	9248	conv2d_16[0][0]
conv2d_18 (Conv2D)	(None,	96,	128	, 32)	9248	conv2d_17[0][0]
conv2d_19 (Conv2D)	(None,	96,	128	, 23)	759	conv2d_18[0][0]

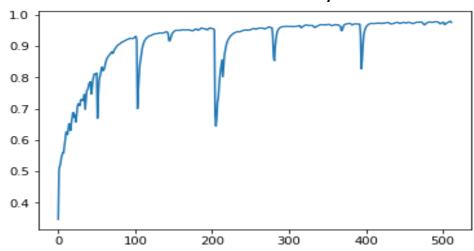
Total params: 8,640,471 Trainable params: 8,640,471 Non-trainable params: 0 [10]:
 unet.summary()

Layer (type)	Output	Shape	Param #	Connected to
input_1 (InputLayer)	[(None	, 96, 128, 3)]	0	
conv2d (Conv2D)	(None,	96, 128, 32)	896	input_1[0][0]
conv2d_1 (Conv2D)	(None,	96, 128, 32)	9248	conv2d[0][0]
max_pooling2d (MaxPooling2D)	(None,	48, 64, 32)	0	conv2d_1[0][0]
conv2d_2 (Conv2D)	(None,	48, 64, 64)	18496	max_pooling2d[0][0]
conv2d_3 (Conv2D)	(None,	48, 64, 64)	36928	conv2d_2[0][0]
max_pooling2d_1 (MaxPooling2D)	(None,	24, 32, 64)	0	conv2d_3[0][0]
conv2d_4 (Conv2D)	(None,	24, 32, 128)	73856	max_pooling2d_1[0][0]
conv2d_5 (Conv2D)	(None,	24, 32, 128)	147584	conv2d_4[0][0]
max_pooling2d_2 (MaxPooling2D)	(None,	12, 16, 128)	0	conv2d_5[0][0]
conv2d_6 (Conv2D)	(None,	12, 16, 256)	295168	max_pooling2d_2[0][0]
conv2d_7 (Conv2D)	(None,	12, 16, 256)	590080	conv2d_6[0][0]
dropout (Dropout)	(None,	12, 16, 256)	0	conv2d_7[0][0]
max_pooling2d_3 (MaxPooling2D)	(None,	6, 8, 256)	0	dropout[0][0]
conv2d_8 (Conv2D)	(None,	6, 8, 512)	1180160	max_pooling2d_3[0][0]
conv2d_9 (Conv2D)	(None,	6, 8, 512)	2359808	conv2d_8[0][0]
dropout_1 (Dropout)	(None,	6, 8, 512)	0	conv2d_9[0][0]
conv2d_transpose (Conv2DTranspo	(None,	12, 16, 256)	1179904	dropout_1[0][0]

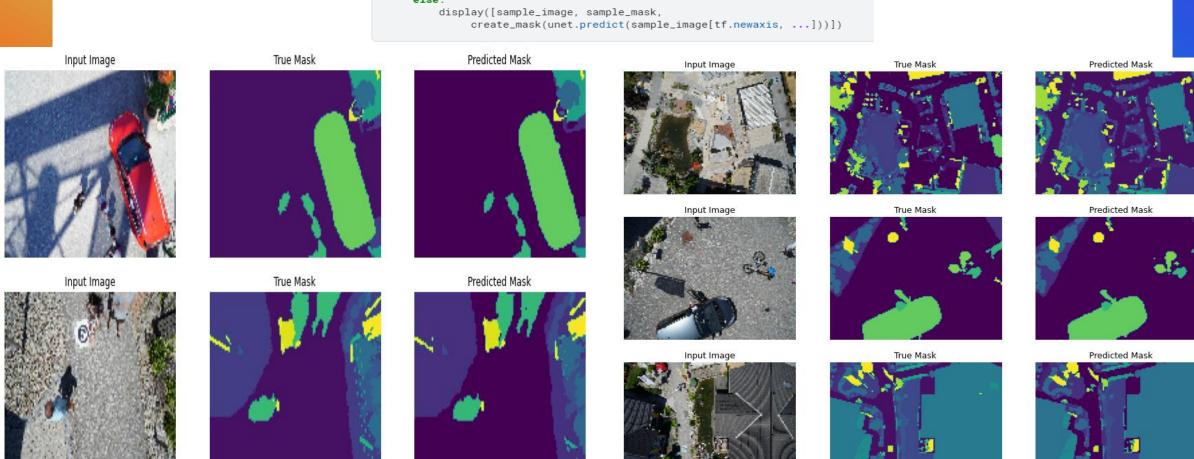
Model Compilation

- + After we finished building the UNet model, we compile this model.
- + For this experiment, we uses
 Adam optimizer, crossentropy
 as loss function and accuracy
 as the metric to measure the
 performance.

```
unet.compile(optimizer='adam',
               loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=True),
               metrics=['accuracy'])
 + Code
            + Markdown
 EPOCHS = 512
 VAL SUBSPLITS = 5
 BUFFER_SIZE = 500
 BATCH_SIZE = 16
 processed_image_ds.batch(BATCH_SIZE)
 train_dataset = processed_image_ds.cache().shuffle(BUFFER_SIZE).batch(BATCH_SIZE)
 print(processed_image_ds.element_spec)
 model_history = unet.fit(train_dataset, epochs=EPOCHS)
(TensorSpec(shape=(96, 128, 3), dtype=tf.float32, name=None), TensorSpec(shape=(96, 128, 1), dtype=tf.uint8, name=None))
Epoch 1/512
                                     Model Accuracy
```



Results



Conclusion

Deep Learning is an continuously-growing and a relatively new concept, the vast amount of resources can be a touch overwhelming for those either looking to get into the field, or those already engraved in it. A good way of cooping is to get a good general knowledge of machine learning and then find a good structured path to follow (be a project or research).



