

# Speech Command Recognition Using Deep Learning

**Hafiz Harun Jamil**  
**AD20216005**  
**Pattern Recognition**

# Introduction

- This presentation shows you how to train a deep learning model that detects the presence of speech commands in audio.
- The example used in this presentation uses the Speech Commands Dataset [1] to train a convolutional neural network to recognize a given set of commands.
- To train a network from scratch, you must first download the data set, or one can load a pretrained network.
- This trained model
  - *Recognize Commands with a Pre-Trained Network*
  - *Detect Commands Using Streaming Audio from Microphone*

# *Detect Commands Using Streaming Audio from Microphone*

- Pre-trained speech recognition network is used to identify speech commands.
- To load the pre-trained network `load('commandNet.mat')` is used.

```
load('commandNet.mat')
```

- The network is trained to recognize the following speech commands:

- "yes"
- "no"
- "up"
- "down"
- "left"
- "right"
- "on"
- "off"
- "stop"
- "go"

- Load a short speech signal where a person says “any of the above mentioned words”.

```
[x,fs] = audioread('stop_command.flac');
```

# Define Neural Network Architecture

- Create a simple network architecture as an array of layers.
- Use convolutional and batch normalization layers, and **downsample** the feature maps "spatially" (that is, in time and frequency) using max pooling layers.
- Add a final max pooling layer that pools the input feature map globally over time.
- This approximate time-translation invariance in the input spectrograms, allowing the network to perform the same classification independent of the exact position of the speech in time.
- Global pooling also significantly reduces the number of parameters in the final fully connected layer.
- To reduce the possibility of the network memorizing specific features of the training data, add a small amount of dropout to the input to the last fully connected layer.

# Cont ...

- The network is small, as it has only five convolutional layers with few filters.
- **numF** controls the number of filters in the convolutional layers.
- To increase the accuracy of the network, try increasing the network depth by adding identical blocks of **convolutional**, **batch normalization**, and **ReLU layers**.
- You can also try increasing the number of convolutional filters by increasing numF.
- Use a weighted cross entropy classification loss.
- **weightedClassificationLayer**(classWeights) creates a custom classification layer that calculates the cross entropy loss with observations weighted by **classWeights**.
- Specify the class weights in the same order as the classes appear in categories(YTrain).
- To give each class equal total weight in the loss, use class weights that are inversely proportional to the number of training examples in each class.
- When using the Adam optimizer to train the network, the training algorithm is independent of the overall normalization of the class weights.

```
classWeights = 1./countcats(YTrain);
classWeights = classWeights'/mean(classWeights);
numClasses = numel(categories(YTrain));

timePoolSize = ceil(numHops/8);

dropoutProb = 0.2;
numF = 12;
layers = [
    imageInputLayer([numHops numBands])

    convolution2dLayer(3,numF,'Padding','same')
    batchNormalizationLayer
    reluLayer

    maxPooling2dLayer(3,'Stride',2,'Padding','same')

    convolution2dLayer(3,2*numF,'Padding','same')
    batchNormalizationLayer
    reluLayer

    maxPooling2dLayer(3,'Stride',2,'Padding','same')

    convolution2dLayer(3,4*numF,'Padding','same')
    batchNormalizationLayer
    reluLayer

    maxPooling2dLayer(3,'Stride',2,'Padding','same')

    convolution2dLayer(3,4*numF,'Padding','same')
    batchNormalizationLayer
    reluLayer
    convolution2dLayer(3,4*numF,'Padding','same')
    batchNormalizationLayer
    reluLayer

    maxPooling2dLayer([timePoolSize,1])

    dropoutLayer(dropoutProb)
    fullyConnectedLayer(numClasses)
    softmaxLayer
    weightedClassificationLayer(classWeights)];
```

# Train Network

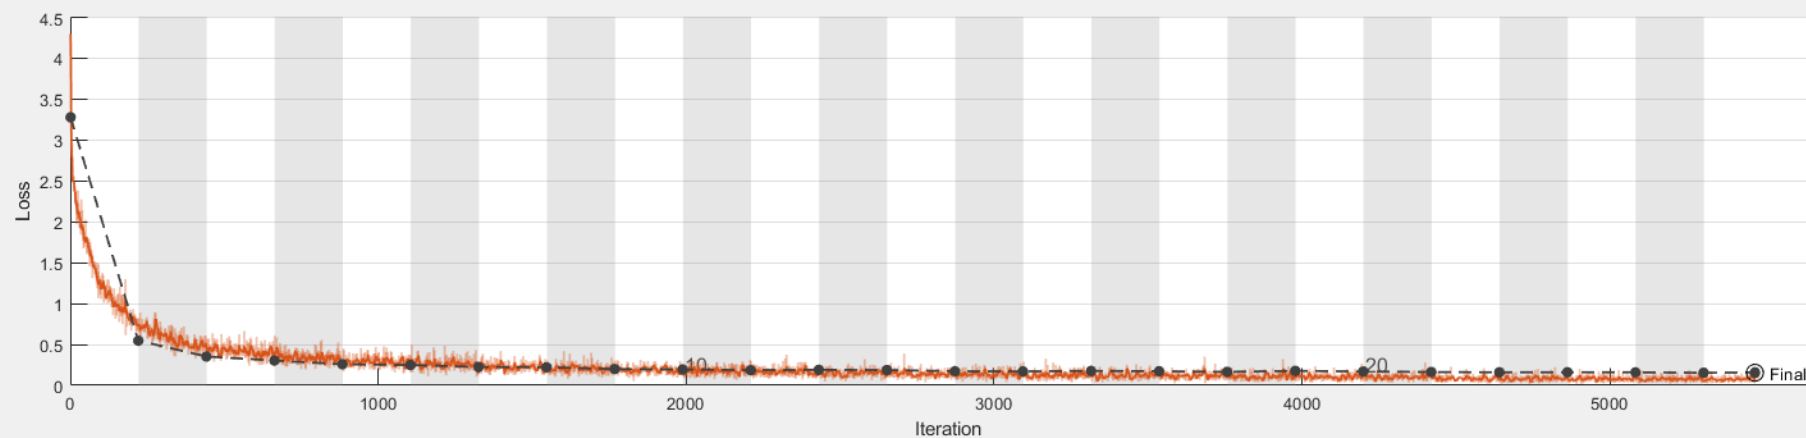
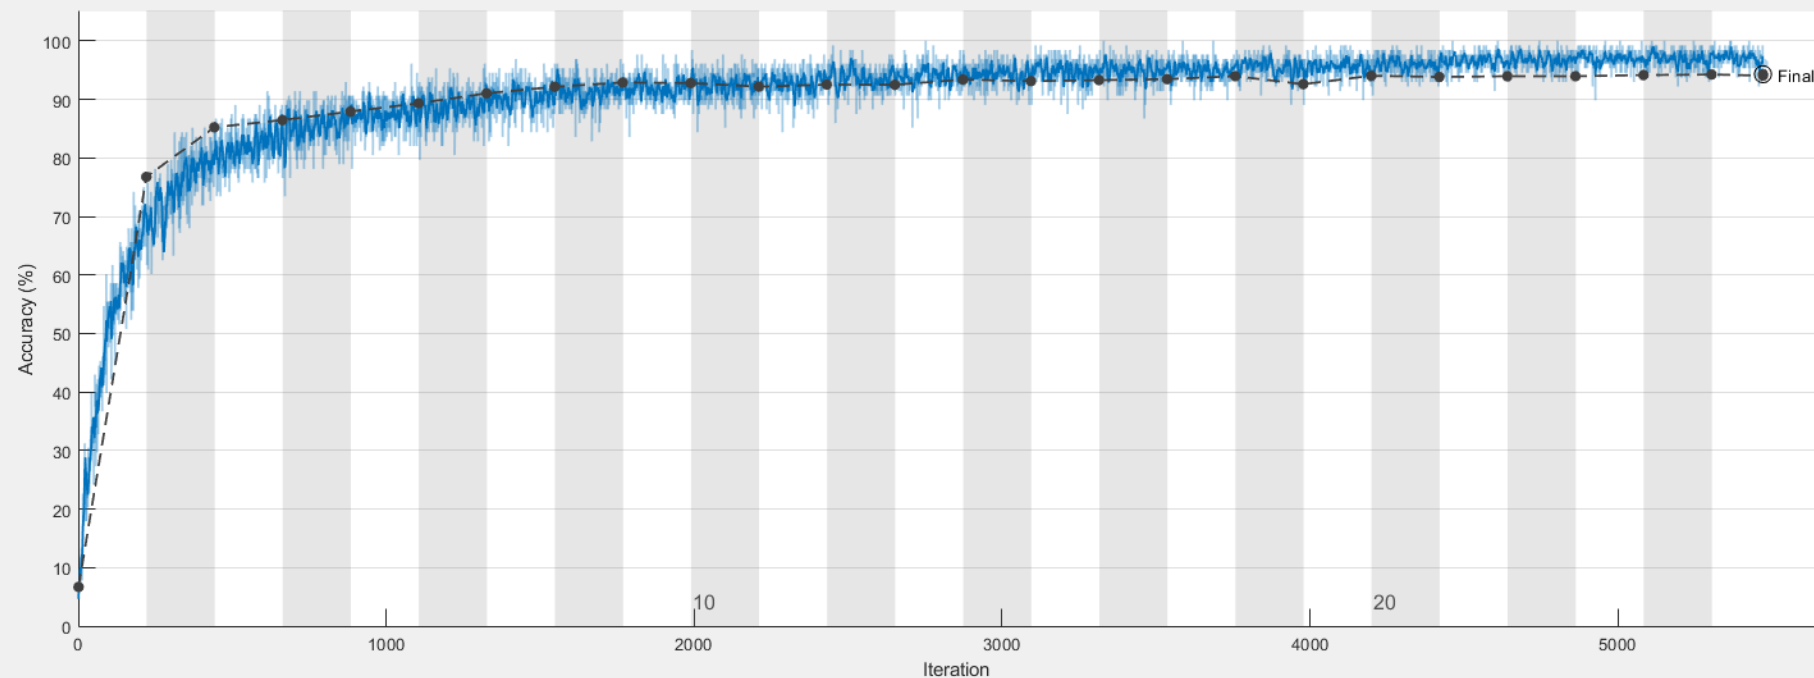
- Specify the training options. Use the Adam optimizer with a mini-batch size of 128. Train for 25 epochs and reduce the learning rate by a factor of 10 after 20 epochs.

```
miniBatchSize = 128;  
validationFrequency = floor(numel(YTrain)/miniBatchSize);  
options = trainingOptions('adam', ...  
    'InitialLearnRate',3e-4, ...  
    'MaxEpochs',25, ...  
    'MiniBatchSize',miniBatchSize, ...  
    'Shuffle','every-epoch', ...  
    'Plots','training-progress', ...  
    'Verbose',false, ...  
    'ValidationData',{XValidation,YValidation}, ...  
    'ValidationFrequency',validationFrequency, ...  
    'LearnRateSchedule','piecewise', ...  
    'LearnRateDropFactor',0.1, ...  
    'LearnRateDropPeriod',20);
```

Train the network by using the command `trainedNet = trainNetwork(XTrain,YTrain,layers,options);`

# Train Network

Training Progress (19-Oct-2021 00:48:31)



## Results

Validation accuracy: 94.29%  
Training finished: Stopped manually

## Training Time

Start time: 19-Oct-2021 00:48:31  
Elapsed time: 5 min 22 sec

## Training Cycle

Epoch: 25 of 25  
Iteration: 5471 of 5525  
Iterations per epoch: 221  
Maximum iterations: 5525

## Validation

Frequency: 221 iterations

## Other Information

Hardware resource: Single GPU  
Learning rate schedule: Piecewise  
Learning rate: 3e-05

[i Learn more](#)

## Accuracy

— Training (smoothed)  
— Training  
- - Validation

## Loss

— Training (smoothed)  
— Training  
- - Validation

# Evaluate Trained Network

- Calculate the final accuracy of the network on the training set (without data augmentation) and validation set.
- The network is very accurate on this data set. However, the training, validation, and test data all have similar distributions that do not necessarily reflect real-world environments.

```
if reduceDataset
    load('commandNet.mat','trainedNet');
end
YValPred = classify(trainedNet,XValidation);
validationError = mean(YValPred ~= YValidation);
YTrainPred = classify(trainedNet,XTrain);
trainError = mean(YTrainPred ~= YTrain);
disp("Training error: " + trainError*100 + "%")
disp("Validation error: " + validationError*100 + "%")
```

### Confusion Matrix for Validation Data

True Class	yes	252	3		1		1			1	3	96.6%	3.4%		
	no		255		5	1				4	5	94.4%	5.6%		
	up		1	244			1	6			5	3	93.8%	6.2%	
	down		10		247		1			3	3		93.6%	6.4%	
	left	3	4			237	1			1		1	96.0%	4.0%	
	right		2	1		3	246				4		96.1%	3.9%	
	on			4				238	6		5	4	92.6%	7.4%	
	off			10				3	242		1		94.5%	5.5%	
	stop			2	1	1			2	234	1	3	2	95.1%	4.9%
	go		13	3	3			1			233	3	4	89.6%	10.4%
	unknown	2	8	6	11	2	10	15	4	5	18	789	1	90.6%	9.4%
	background												600	100.0%	
98.1% 86.1% 90.4% 92.2% 97.1% 95.7% 91.5% 93.1% 97.9% 89.6% 96.3% 97.1%															
1.9% 13.9% 9.6% 7.8% 2.9% 4.3% 8.5% 6.9% 2.1% 10.4% 3.7% 2.9%															
yes no up down left right on off stop go unknown background															
Predicted Class															



# Cont...

- When working on applications with constrained hardware resources such as mobile applications, consider the limitations on available memory and computational resources.
- Compute the total size of the network in kilobytes and test its prediction speed when using a CPU.
- The prediction time is the time for classifying a single input image.
- If you input multiple images to the network, these can be classified simultaneously, leading to shorter prediction times per image.
- When classifying streaming audio, however, the single-image prediction time is the most relevant.

```
info = whos('trainedNet');  
disp("Network size: " + info.bytes/1024 + " kB")  
  
for i = 1:100  
    x = randn([numHops,numBands]);  
    tic  
    [YPredicted,probs] = classify(trainedNet,x,"ExecutionEnvironment",'cpu');  
    time(i) = toc;  
end  
disp("Single-image prediction time on CPU: " + mean(time(11:end))*1000 + " ms")
```

Network size: 286.7402 kB

Single-image prediction time on CPU: 2.495 ms

# Pre-trained network takes auditory-based spectrograms

- Listen to the command by using this command.

```
sound(x,fs)
```

- The pre-trained network takes auditory-based spectrograms as inputs.
- First convert the speech waveform to an auditory-based spectrogram by using the function “**extractAuditoryFeature**”.

```
auditorySpect = helperExtractAuditoryFeatures(x,fs);
```

- The next step is to classify the command based on its auditory spectrogram by using the command

```
command = classify(trainedNet,auditorySpect)
```

# Unknown Auditory Spectrogram

- For the case of unknown auditory spectrogram, the network is trained to classify words not belonging to this set as "unknown".
- For example:
- Classify a word ("any") that was not included in the list (slide:3) of command to identify.
- Load the speech signal and listen to it by using command

```
x = audioread('play_command.flac');  
sound(x,fs)
```

- Compute the auditory spectrogram

```
auditorySpect = helperExtractAuditoryFeatures(x,fs);
```

- Classify the signal.

```
command = classify(trainedNet,auditorySpect)
```

```
command =
```

```
categorical
```

```
unknown
```

# Background Noise

- It is important to consider the noise factor in speech recognition, for that purpose the network is trained to classify background noise as "background".
- Create a one-second signal consisting of random noise by using command

```
x = pinknoise(16e3);
```

- Compute the auditory spectrogram

```
auditorySpect = helperExtractAuditoryFeatures(x,fs);
```

- Classify the background noise

```
command = classify(trainedNet,auditorySpect)
```

```
command =
```

```
categorical
```

```
background
```

# Detect Commands Using Streaming Audio from Microphone

- To test pre-trained command detection network on streaming audio from your microphone.
- Try saying one of the commands, from the list in slide 3.
- Then, try saying one of the unknown words such as *Marvin*, *Sheila*, *bed*, *house*, *cat*, *bird*, or any number from zero to nine.
- Specify the classification rate in Hz and create an audio device reader that can read audio from your microphone by using the command

```
classificationRate = 20;  
adr = audioDeviceReader('SampleRate',fs,'SamplesPerFrame',floor(fs/classificationRate));
```

# Buffer Initialize

- Initialize a buffer for the audio.
- Extract the classification labels of the network.
- Initialize buffers of half a second for the labels and classification probabilities of the streaming audio.
- Use these buffers to compare the classification results over a longer period of time and by that build 'agreement' over when a command is detected.
- Specify thresholds for the decision logic.

```
audioBuffer = dsp.AsyncBuffer(fs);  
  
labels = trainedNet.Layers(end).Classes;  
YBuffer(1:classificationRate/2) = categorical("background");  
  
probBuffer = zeros([numel(labels),classificationRate/2]);  
  
countThreshold = ceil(classificationRate*0.2);  
probThreshold = 0.7;
```

# Cont ...

- Create a figure and detect commands as long as the created figure exists.
- To run the loop indefinitely, set timeLimit to Inf.
- To stop the live detection, simply close the figure.

```
h = figure('Units','normalized','Position',[0.2 0.1 0.6 0.8]);

timeLimit = 20;

tic
while ishandle(h) && toc < timeLimit
```

```
x = adr();
write(audioBuffer,x);
y = read(audioBuffer,fs,fs-adr.SamplesPerFrame);

spec = helperExtractAuditoryFeatures(y,fs);
```

```
[YPredicted,probs] = classify(trainedNet,spec,'ExecutionEnvironment','cpu');
YBuffer = [YBuffer(2:end),YPredicted];
probBuffer = [probBuffer(:,2:end),probs(:)];
```

```
subplot(2,1,1)
plot(y)
axis tight
ylim([-1,1])

subplot(2,1,2)
pcolor(spec')
caxis([-4 2.6445])
shading flat
```

```
[YMode,count] = mode(YBuffer);

maxProb = max(probBuffer(labels == YMode,:));
subplot(2,1,1)
if YMode == "background" || count < countThreshold || maxProb < probThreshold
    title(" ")
else
    title(string(YMode),'FontSize',20)
end

drawnow
end
```

# Load Speech Commands Data Set

- This example uses the Google Speech Commands Dataset [1].
- Download the dataset and untar the downloaded file.
- Set PathToDatabase to the location of the data.

.	2,055,161,8...	? File folder	4/12/2018 2:45 ...
. \zero	75,377,540	? File folder	7/28/2017 5:10 ...
. \stop	75,172,872	? File folder	7/28/2017 5:10 ...
. \seven	75,165,002	? File folder	7/28/2017 5:10 ...
. \six	75,133,166	? File folder	7/28/2017 5:10 ...
. \yes	75,056,272	? File folder	7/28/2017 5:10 ...
. \four	74,925,542	? File folder	7/28/2017 5:10 ...
. \two	74,830,770	? File folder	7/28/2017 5:10 ...
. \nine	74,806,158	? File folder	7/28/2017 5:10 ...
. \right	74,754,036	? File folder	7/28/2017 5:10 ...
. \no	74,722,134	? File folder	7/28/2017 5:10 ...
. \down	74,611,672	? File folder	7/28/2017 5:10 ...
. \five	74,569,366	? File folder	7/28/2017 5:10 ...
. \one	74,559,656	? File folder	7/28/2017 5:10 ...
. \up	74,510,102	? File folder	7/28/2017 5:10 ...
. \off	74,506,438	? File folder	7/28/2017 5:10 ...
. \left	74,506,080	? File folder	7/28/2017 5:10 ...
. \go	74,496,148	? File folder	7/28/2017 5:10 ...
. \on	74,472,528	? File folder	7/28/2017 5:10 ...
. \three	74,423,116	? File folder	7/28/2017 5:10 ...
. \eight	74,182,774	? File folder	7/28/2017 5:10 ...
. \house	55,159,508	? File folder	7/28/2017 5:10 ...
. \marvin	55,155,270	? File folder	7/28/2017 5:10 ...
. \dog	54,982,048	? File folder	7/28/2017 5:10 ...
. \happy	54,922,478	? File folder	7/28/2017 5:10 ...
. \sheila	54,888,718	? File folder	7/28/2017 5:10 ...
. \wow	54,836,460	? File folder	7/28/2017 5:10 ...
. \tree	54,499,448	? File folder	7/28/2017 5:10 ...
. \cat	54,476,952	? File folder	7/28/2017 5:10 ...
. \bird	54,409,302	? File folder	7/28/2017 5:10 ...
. \bed	53,877,116	? File folder	7/28/2017 5:10 ...
. \_background_noise_	12,782,147	? File folder	7/28/2017 5:11 ...



# Create Training Datastore

- Create an `audioDatastore` (Audio Toolbox) that points to the training data set.

```
ads = audioDatastore(fullfile(dataFolder, 'train'), ...  
    'IncludeSubfolders',true, ...  
    'FileExtensions','.wav', ...  
    'LabelSource','foldernames')
```

## Choose Words to Recognize

- Specify the words that you want your model to recognize as commands.
- Label all words that are not commands as unknown.
- Labeling words that are not commands as unknown creates a group of words that approximates the distribution of all words other than the commands.
- The network uses this group to learn the difference between commands and all other words.
- **To reduce the class imbalance between the known and unknown words and speed up processing, only include a fraction of the unknown words in the training set.**

```
commands = categorical(["yes","no","up","down","left","right","on","off","stop","go"]);  
  
isCommand = ismember(ads.Labels,commands);  
isUnknown = ~isCommand;  
  
includeFraction = 0.2;  
mask = rand(numel(ads.Labels),1) < includeFraction;  
isUnknown = isUnknown & mask;  
ads.Labels(isUnknown) = categorical("unknown");  
  
adsTrain = subset(ads,isCommand|isUnknown);  
countEachLabel(adsTrain)
```

ans =

11×2 table

Label	Count
down	1842
go	1861
left	1839
no	1853
off	1839
on	1864
right	1852
stop	1885
unknown	6483
up	1843
yes	1860

# Create Validation Datastore

- Create an audioDatastore (Audio Toolbox) that points to the validation data set. Follow the same steps used to create the training datastore.

```
ads = audioDatastore(fullfile(dataFolder, 'validation'), ...  
    'IncludeSubfolders',true, ...  
    'FileExtensions','.wav', ...  
    'LabelSource','foldernames')  
  
isCommand = ismember(ads.Labels,commands);  
isUnknown = ~isCommand;  
  
includeFraction = 0.2;  
mask = rand(numel(ads.Labels),1) < includeFraction;  
isUnknown = isUnknown & mask;  
ads.Labels(isUnknown) = categorical("unknown");  
  
adsValidation = subset(ads,isCommand|isUnknown);  
countEachLabel(adsValidation)
```

```
ads =  
  
audioDatastore with properties:  
  
    Files: {  
        '...\AppData\Local\Temp\google_speech\validation\bed\026290a7_nohash_0.wav';  
        '...\AppData\Local\Temp\google_speech\validation\bed\060cd039_nohash_0.wav';  
        '...\AppData\Local\Temp\google_speech\validation\bed\060cd039_nohash_1.wav'  
        ... and 6795 more  
    }  
    Folders: {  
        'C:\Users\jibrahim\AppData\Local\Temp\google_speech\validation'  
    }  
    Labels: [bed; bed; bed ... and 6795 more categorical]  
AlternateFileSystemRoots: {}  
    OutputDataType: 'double'  
SupportedOutputFormats: ["wav"    "flac"    "ogg"    "mp4"    "m4a"]  
    DefaultOutputFormat: "wav"
```

```
ans =  
  
11x2 table  
  
    Label    Count  
    _____  _____  
    down      264  
    go        260  
    left      247  
    no        270  
    off       256  
    on        257  
    right     256  
    stop      246  
    unknown   850  
    up        260  
    yes       261
```

# Compute Auditory Spectrograms

- To prepare the data for efficient training of a convolutional neural network, convert the speech waveforms to auditory-based spectrograms.
- Define the parameters of the feature extraction.
- **segmentDuration** is the duration of each speech clip (in seconds).
- **frameDuration** is the duration of each frame for spectrum calculation.
- **hopDuration** is the time step between each spectrum.
- **numBands** is the number of filters in the auditory spectrogram.

# Cont...

- Create an **audioFeatureExtractor** (Audio Toolbox) object to perform the feature extraction.

```
fs = 16e3; % Known sample rate of the data set.

segmentDuration = 1;
frameDuration = 0.025;
hopDuration = 0.010;

segmentSamples = round(segmentDuration*fs);
frameSamples = round(frameDuration*fs);
hopSamples = round(hopDuration*fs);
overlapSamples = frameSamples - hopSamples;

FFTLength = 512;
numBands = 50;

afe = audioFeatureExtractor( ...
    'SampleRate',fs, ...
    'FFTLength',FFTLength, ...
    'Window',hann(frameSamples,'periodic'), ...
    'OverlapLength',overlapSamples, ...
    'barkSpectrum',true);
setExtractorParams(afe,'barkSpectrum','NumBands',numBands,'WindowNormalization',false);
```

## Note:

- Read a file from the dataset.
- Training a convolutional neural network requires input to be a consistent size.
- Some files in the data set are less than 1 second long.
- Apply zero-padding to the front and back of the audio signal so that it is of length `segmentSamples`.

```
x = read(adsTrain);

numSamples = size(x,1);

numToPadFront = floor( (segmentSamples - numSamples)/2 );
numToPadBack = ceil( (segmentSamples - numSamples)/2 );

xPadded = [zeros(numToPadFront,1,'like',x);x;zeros(numToPadBack,1,'like',x)];
```

# Cont ...

- To extract audio features, call `extract`. The output is a Bark spectrum with time across rows.

```
features = extract(afe,xPadded);  
[numHops,numFeatures] = size(features)
```

- To speed up processing, you can distribute the feature extraction across multiple workers using `parfor`.
- First, determine the number of partitions for the dataset.
- If you do not have [Parallel Computing Toolbox™](#), use a single partition.

```
if ~isempty(ver('parallel')) && ~reduceDataset  
    pool = gcp;  
    numPar = numpartitions(adsTrain,pool);  
else  
    numPar = 1;  
end
```

- For each partition, read from the datastore, zero-pad the signal, and then extract the features.

```
parfor ii = 1:numPar  
    subds = partition(adsTrain,numPar,ii);  
    XTrain = zeros(numHops,numBands,1,numel(subds.Files));  
    for idx = 1:numel(subds.Files)  
        x = read(subds);  
        xPadded = [zeros(floor((segmentSamples-size(x,1))/2),1);x;zeros(ceil((segmentSamples-size(x,1))/2),1)];  
        XTrain(:,:,idx) = extract(afe,xPadded);  
    end  
    XTrainC{ii} = XTrain;  
end
```

# Cont ...

- Convert the output to a 4-dimensional array with auditory spectrograms along the fourth dimension.

```
XTrain = cat(4,XTrainC{:});  
  
[numHops,numBands,numChannels,numSpec] = size(XTrain)
```

- To obtain data with a smoother distribution, take the logarithm of the spectrograms using a small offset.

```
epsil = 1e-6;  
XTrain = log10(XTrain + epsil);
```

- Perform the feature extraction steps described above to the validation set.

```
if ~isempty(ver('parallel'))  
    pool = gcp;  
    numPar = numpartitions(adsValidation,pool);  
else  
    numPar = 1;  
end  
parfor ii = 1:numPar  
    subds = partition(adsValidation,numPar,ii);  
    XValidation = zeros(numHops,numBands,1,numel(subds.Files));  
    for idx = 1:numel(subds.Files)  
        x = read(subds);  
        xPadded = [zeros(floor((segmentSamples-size(x,1))/2),1);x;zeros(ceil((segmentSamples-size(x,1))/2),1)];  
        XValidation(:, :, :, idx) = extract(afe,xPadded);  
    end  
    XValidationC{ii} = XValidation;  
end  
XValidation = cat(4,XValidationC{:});  
XValidation = log10(XValidation + epsil);
```

# Cont ...

- Isolate the train and validation labels.
- Remove empty categories.

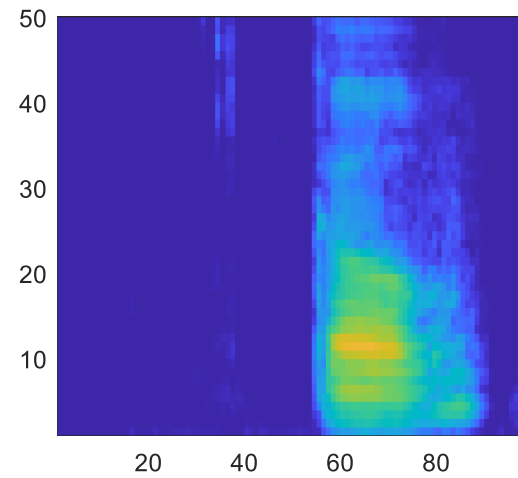
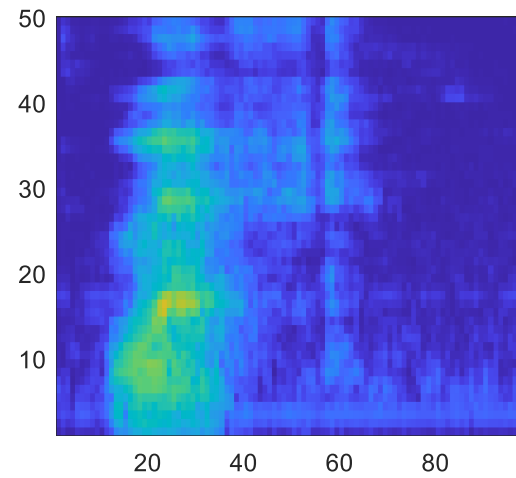
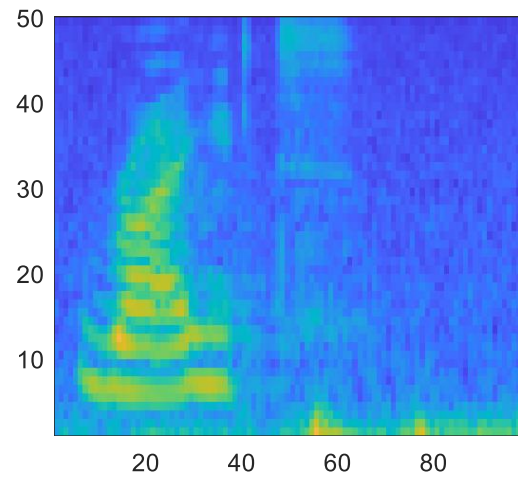
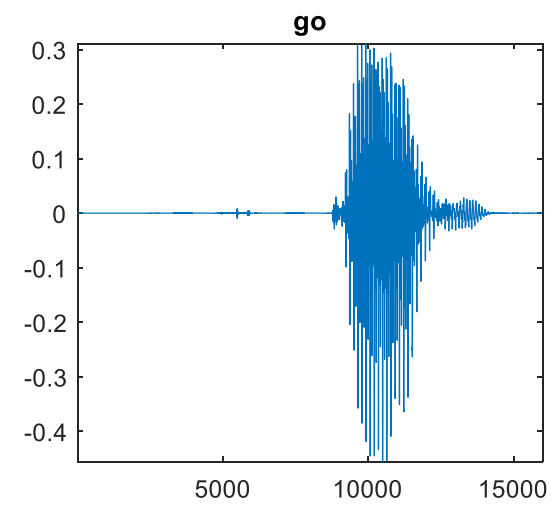
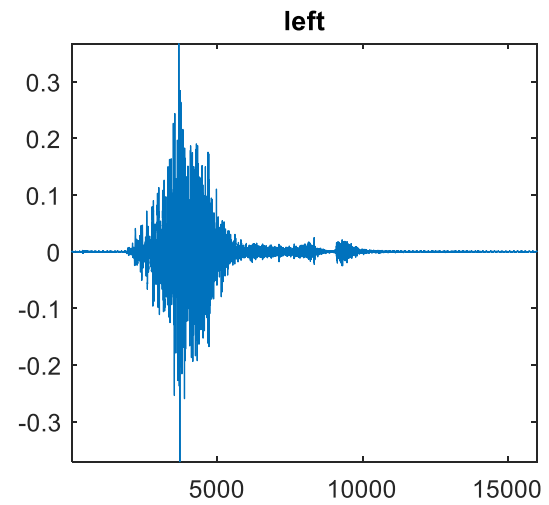
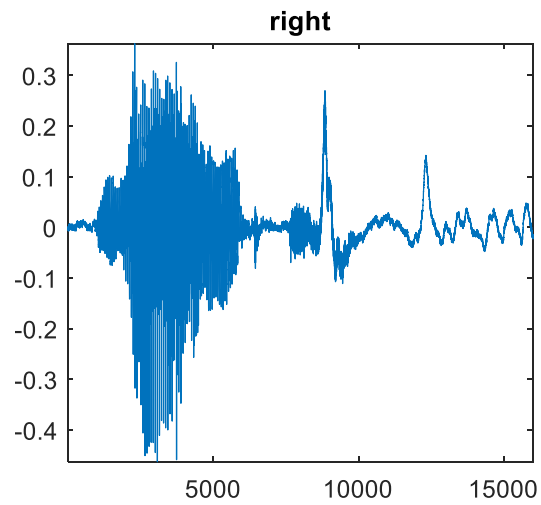
```
YTrain = removecats(adsTrain.Labels);  
YValidation = removecats(adsValidation.Labels);
```

## Visualize Data

- Plot the waveforms and auditory spectrograms of a few training samples. Play the corresponding audio clips.

```
specMin = min(XTrain,[],'all');  
specMax = max(XTrain,[],'all');  
idx = randperm(numel(adsTrain.Files),3);  
figure('Units','normalized','Position',[0.2 0.2 0.6 0.6]);  
for i = 1:3  
    [x,fs] = audioread(adsTrain.Files{idx(i)});  
    subplot(2,3,i)  
    plot(x)  
    axis tight  
    title(string(adsTrain.Labels{idx(i)}))  
  
    subplot(2,3,i+3)  
    spect = (XTrain(:, :, 1, idx(i)))';  
    pcolor(spect)  
    caxis([specMin specMax])  
    shading flat  
  
    sound(x,fs)  
    pause(2)  
end
```

# Cont...





# Add Background Noise Data

- The network must be able not only to recognize different spoken words but also to detect if the input contains silence or background noise.
- Use the audio files in the `_background_` folder to create samples of one-second clips of background noise.
- Create an equal number of background clips from each background noise file. You can also create your own recordings of background noise and add them to the `_background_` folder.
- Before calculating the spectrograms, the function rescales each audio clip with a factor sampled from a log-uniform distribution in the range given by `volumeRange`.

```
adsBkg = audioDatastore(fullfile(dataFolder, 'background'))
numBkgClips = 4000;
if reduceDataset
    numBkgClips = numBkgClips/20;
end
volumeRange = log10([1e-4,1]);

numBkgFiles = numel(adsBkg.Files);
numClipsPerFile = histcounts(1:numBkgClips, linspace(1,numBkgClips,numBkgFiles+1));
Xbkg = zeros(size(XTrain,1),size(XTrain,2),1,numBkgClips,'single');
bkgAll = readall(adsBkg);
ind = 1;

for count = 1:numBkgFiles
    bkg = bkgAll{count};
    idxStart = randi(numel(bkg)-fs,numClipsPerFile(count),1);
    idxEnd = idxStart+fs-1;
    gain = 10.^((volumeRange(2)-volumeRange(1))*rand(numClipsPerFile(count),1) + volumeRange(1));
    for j = 1:numClipsPerFile(count)

        x = bkg(idxStart(j):idxEnd(j))*gain(j);

        x = max(min(x,1),-1);

        Xbkg(:,:,,ind) = extract(afe,x);

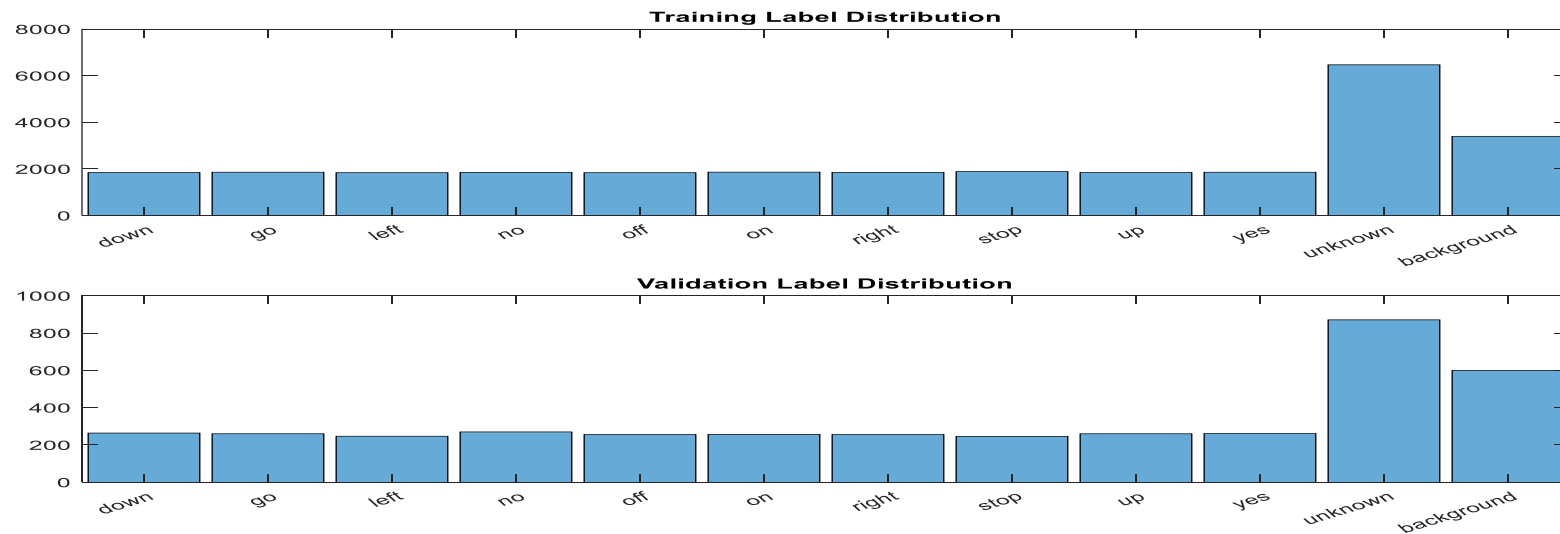
        if mod(ind,1000)==0
            disp("Processed " + string(ind) + " background clips out of " + string(numBkgClips))
        end
        ind = ind + 1;
    end
end
Xbkg = log10(Xbkg + epsil);
```

# Cont...

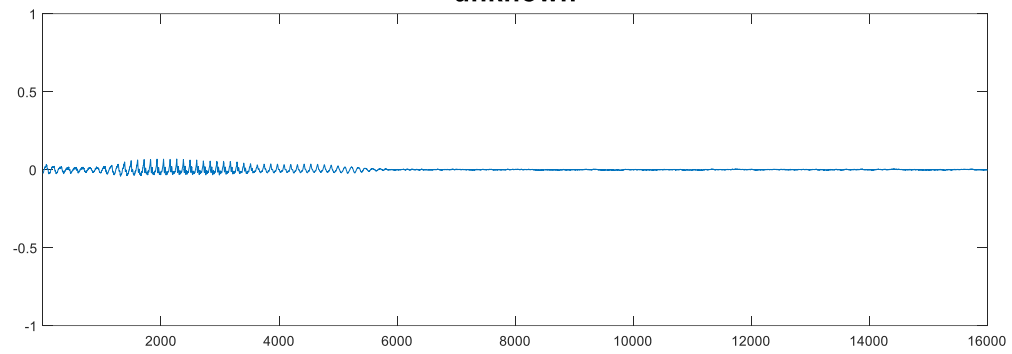
- Split the spectrograms of background noise between the training, validation, and test sets.
- Because the `_background_noise_` folder contains only about five and a half minutes of background noise, the background samples in the different data sets are highly correlated.
- To increase the variation in the background noise, you can create your own background files and add them to the folder.
- To increase the robustness of the network to noise, you can also try mixing background noise into the speech files.

```
numTrainBkg = floor(0.85*numBkgClips);  
numValidationBkg = floor(0.15*numBkgClips);  
  
XTrain(:,:,end+1:end+numTrainBkg) = Xbkg(:,:,1:numTrainBkg);  
YTrain(end+1:end+numTrainBkg) = "background";  
  
XValidation(:,:,end+1:end+numValidationBkg) = Xbkg(:,:,numTrainBkg+1:end);  
YValidation(end+1:end+numValidationBkg) = "background";
```

- Plot the distribution of the different class labels in the training and validation sets.



unknown



up

