

Face Mask Detection - VGG16

AM20216801

JANG CHEOL
HEE

CONTENTS

01

VGG16

02

Library

03

Data set

04


Import model

05

Model


06

Conclusion




VGG16

- ✓ VGGNet is a model developed by VGG, a research team at Oxford University, and is a model that won the 2014 ImageNet image recognition contest.
- ✓ VGGNet represents a model with 16 or 19 layers.
- ✓ VGG16 consists of 13 convolutional layers and 3 fully connected layers.



VGG16 구조



VGG16 구조



Library

- numpy
- keras
- shutil
- os
- random
- matplotlib

```
#importing the necessary libraries
import numpy as np
import keras
import shutil
import os
import random
import matplotlib.pyplot as plt
%matplotlib inline

from keras.preprocessing.image import ImageDataGenerator

[ ] from google.colab import drive
    drive.mount('/content/drive')
```



Data set

- 209 mask images and 132 no mask images were used.
- Selected photos randomly and classified train data and test data.
- Used ImageDataGenerator to preprocess by dividing by 255 and augmentation the image.
- Data using flow_from_directory
Batch, Target Size Class Mode Adjustment
“Categorical”.

```
[ ] #assigning image paths to variables
mask_data = "/content/sample_data/Mask/"
no_mask_data = "/content/sample_data/No Mask/"

[ ] total_mask_images = os.listdir(mask_data)
print("no of mask images:: {}".format(len(total_mask_images)))
total_nonmask_images = os.listdir(no_mask_data)
print("no of non-mask images:: {}".format(len(total_nonmask_images)))

no of mask images:: 209
no of non-mask images:: 132
```

```
▶ for images in random.sample(total_mask_images,100):
    shutil.copy(mask_data+images, '/content/train/mask')
for images in random.sample(total_mask_images,30):
    shutil.copy(mask_data+images, '/content/test/mask')
for images in random.sample(total_nonmask_images,100):
    shutil.copy(no_mask_data+images, '/content/train/no mask')
for images in random.sample(total_nonmask_images,30):
    shutil.copy(no_mask_data+images, '/content/test/no mask')

[ ] train_batch = ImageDataGenerator(rescale=1./255, zoom_range=0.2, horizontal_flip=True, vertical_flip=True, shear_range=0.2).#
    flow_from_directory('./train', target_size=(224,224), batch_size=32, class_mode = 'categorical')
test_batch = ImageDataGenerator(rescale=1./255).#
    flow_from_directory('./test', target_size = (224,224), batch_size=32, class_mode='categorical')

Found 197 images belonging to 3 classes.
Found 60 images belonging to 3 classes.
```

```
[ ] train_batch.class_indices

{'./ipynb_checkpoints': 0, 'mask': 1, 'no mask': 2}

[ ] class_mask = ['miss', 'mask', 'no mask']
```

Import model

```
[ ] #import vgg16
    from keras.applications.vgg16 import VGG16
```

```
▶ #vgg16 accepts image size (224,224) only
  IMAGE_SIZE = [224,224]
  vgg = VGG16(input_shape=IMAGE_SIZE+[3], weights='imagenet', include_top=False)
```

```
↳ Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16\_weights\_tf\_dim\_ordering\_tf\_kernels\_notop.h5
58892288/58889256 [=====] - 2s 0us/step
58900480/58889256 [=====] - 2s 0us/step
```

```
[ ] for layers in vgg.layers:
    layers.trainable = False
```

```
▶ vgg.summary()
```

```
▶ vgg.summary()
```


Model: "vgg16"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
Total params: 14,714,688		
Trainable params: 0		
Non-trainable params: 14,714,688		

import vgg16 and input_shape = [224,224,3],
weights = "imagenet",
With include_top = False, not to import fully connected layers.

layers.trainable = False to prevent it from being updated.

Total params became non-trainable params.



Model

- Create a flatten layer.
- The output of the dense layer is three activations = “softmax”.
- A flatten layer and a dense layer are added to the existing model

```
[ ] flatten_layer = keras.layers.Flatten()(vgg.output)
prediction_layer = keras.layers.Dense(3, activation='softmax')(flatten_layer)


[ ] model = keras.models.Model(inputs = vgg.input, outputs = prediction_layer)

model.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
flatten (Flatten)	(None, 25088)	0
dense (Dense)	(None, 3)	75267

Total params: 14,789,955
Trainable params: 75,267
Non-trainable params: 14,714,688



Model

```
[ ] model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])

▶ r = model.fit_generator(train_batch, validation_data=test_batch, epochs=5, steps_per_epoch=len(train_batch), validation_steps=len(test_batch))

↳ /usr/local/lib/python3.7/dist-packages/keras/engine/training.py:1972: UserWarning: `Model.fit_generator` is deprecated and will be removed in a future version. Please use `Model.fit`, which supports generators.
  warnings.warn("`Model.fit_generator` is deprecated and ")
/usr/local/lib/python3.7/dist-packages/PIL/TiffImagePlugin.py:788: UserWarning: Corrupt EXIF data. Expecting to read 4 bytes but only got 0.
  warnings.warn(str(msg))
Epoch 1/5
7/7 [=====] - 56s 3s/step - loss: 0.9788 - accuracy: 0.5127 - val_loss: 0.4102 - val_accuracy: 0.8167
Epoch 2/5
7/7 [=====] - 7s 1s/step - loss: 0.4650 - accuracy: 0.7716 - val_loss: 0.1781 - val_accuracy: 0.9333
Epoch 3/5
7/7 [=====] - 7s 1s/step - loss: 0.3088 - accuracy: 0.9036 - val_loss: 0.1741 - val_accuracy: 0.9500
Epoch 4/5
7/7 [=====] - 7s 1s/step - loss: 0.1740 - accuracy: 0.9442 - val_loss: 0.0788 - val_accuracy: 0.9833
Epoch 5/5
7/7 [=====] - 7s 1s/step - loss: 0.1254 - accuracy: 0.9645 - val_loss: 0.0518 - val_accuracy: 0.9833
```

Use optimizer = “adam”.

Use loss = “categorical_crossentropy”.

metrics = [“accuracy”]

- loss : training loss value

- acc : training accuracy

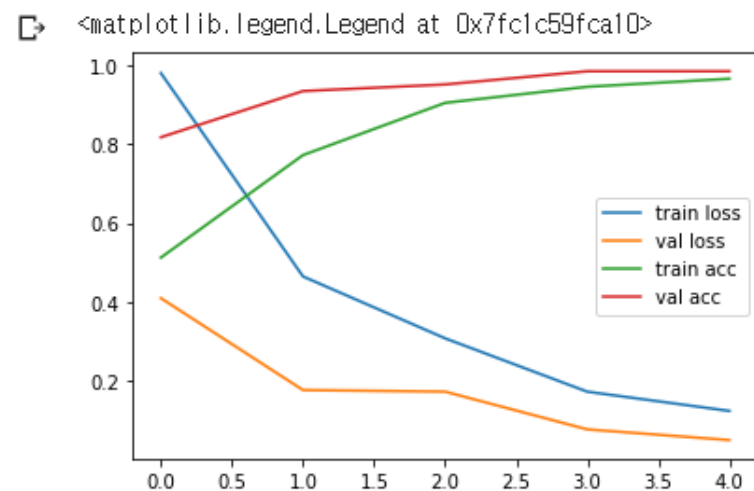
- val_loss : validation loss value

- val_acc : verification accuracy

Result

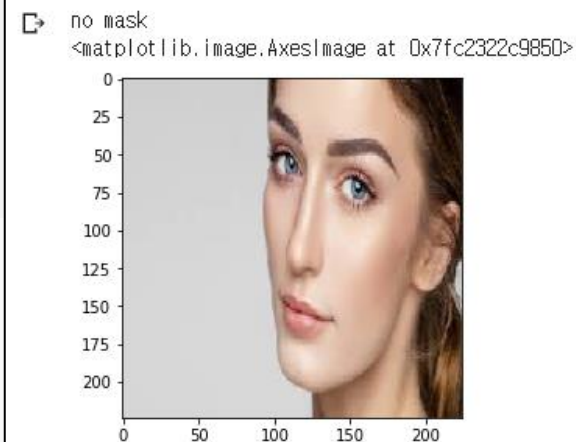
```
plt.plot(r.history['loss'], label = 'train loss')
plt.plot(r.history['val_loss'], label='val loss')
plt.legend()

plt.plot(r.history['accuracy'], label = 'train acc')
plt.plot(r.history['val_accuracy'], label='val acc')
plt.legend()
```



```
from keras.preprocessing import image
from keras.applications.imagenet_utils import preprocess_input
```


```
img = image.load_img('/content/test/no mask/No Mask112.jpg', target_size=(224,224))
x=image.img_to_array(img)
x = np.expand_dims(x,0)
y = preprocess_input(x)
pred = class_mask[np.argmax(model.predict(y))]
print(pred)
plt.imshow(img)
```





Conclusion

- After classifying whether or not a mask was used with vgg16, approach with 0.1254 loss and 0.9645 accuracy at 5 epoch.



Reference

- <https://bskyvision.com/504>
- <https://www.kaggle.com/arkajitmajumder/face-mask-detection/notebook>
- <https://eremo2002.tistory.com/57?category=779320>